

CIT237

Chapter 17:

The Standard Template Library

November 18, 2019

NOTICE to Students: Portions of these lecture slides are
Copyright 2018 Pearson Education, Inc.

We have a license to use this material *only* in the context of this course.
There is NO PERMISSION to share these lecture slides with anyone not taking the course.
There is NO PERMISSION to post these lecture slides on the Internet.

Reminders

- Quiz 6 will be held at the start of class on
Wednesday, November 20.
- The material covered on Quiz 6 will be:
 - Lectures of October 28 through November 13.
 - Chapters 15 and 16.
- Project 3:
Due date is December 2.

The Standard Template Library

- **The Standard Template Library (STL):** an extensive library of generic templates for classes and functions.
- **Categories of Templates:**
 - **Containers:** Class templates for objects that store and organize data
 - **Iterators:** Class templates for objects that behave like pointers, and are used to access the individual data elements in a container
 - **Algorithms:** Function templates that perform various operations on elements of containers

Containers

- **Sequence Containers**
 - Stores data “sequentially”, in a fashion similar to an array: they appear (to the programmer) that they are sequential.
 - Some sequence containers *actually* use contiguous memory addresses, others do not.
- **Associative Containers**
 - Store data in a nonsequential way that can make it faster to locate particular elements in the container.

STL Header Files

Table 17-4 Header Files

Header File	Classes
<code><array></code>	<code>array</code>
<code><deque></code>	<code>deque</code>
<code><forward_list></code>	<code>forward_list</code>
<code><list></code>	<code>list</code>
<code><map></code>	<code>map</code> , <code>multimap</code>
<code><queue></code>	<code>queue</code> , <code>priority_queue</code>
<code><set></code>	<code>set</code> , <code>multiset</code>
<code><stack></code>	<code>stack</code>
<code><unordered_map></code>	<code>unordered_map</code> , <code>unordered_multimap</code>
<code><unordered_set></code>	<code>unordered_set</code> , <code>unordered_multiset</code>
<code><vector></code>	<code>vector</code>

The array Class Template (1)

- An array object works very much like a regular array
- A fixed-size container that holds elements of the same data type.
- array objects have a **size()** member function that returns the number of elements contained in the object.

The array Class Template (2)

- The array class is declared in the <array> header file.
- When defining an array object, you specify the data type of its elements, and the number of elements.
- Examples:

```
array<int, 5> numbers;
```

```
array<string, 4> names;
```

Iterators

- Objects that work like pointers
- Used to access data in STL containers
- Five categories of iterators:

Table 17-6 Categories of Iterators

Iterator Category	Description
Forward	Can only move forward in a container (uses the ++ operator).
Bidirectional	Can move forward or backward in a container (uses the ++ and -- operators).
Random access	Can move forward and backward, and can jump to a specific data element in a container.
Input	Can be used with an input stream to read data from an input device or a file.
Output	Can be used with an output stream to write data to an output device or a file.

Similarities between Pointers and Iterators

	Pointers	Iterators
Use the * and -> operators to dereference	Yes	Yes
Use the = operator to assign to an element	Yes	Yes
Use the == and != operators to compare	Yes	Yes
Use the ++ operator to increment	Yes	Yes
Use the -- operator to decrement	Yes	Yes (bidirectional and random- access iterators)
Use the + operator to move forward a specific number of elements	Yes	Yes
Use the - operator to move backward a specific number of elements	Yes	Yes (bidirectional and random- access iterators)

Defining an Iterator

- To define an iterator, you must know what type of container you will be using it with.
- The general format of an iterator definition:

containerType::iterator iteratorName;

Where *containerType* is the STL container type, and *iteratorName* is the name of the iterator variable that you are defining.

Iterator Examples (1)

- For example, suppose we have defined an `array` object, as follows:

```
array<string, 3> names = {"Sarah", "William", "Alfredo"};
```

We can define an iterator that is compatible with the `array` object as follows:

```
array<string, 3>::iterator it;
```

This defines an iterator named `it`. The iterator can be used with an `array<string, 3>` object.

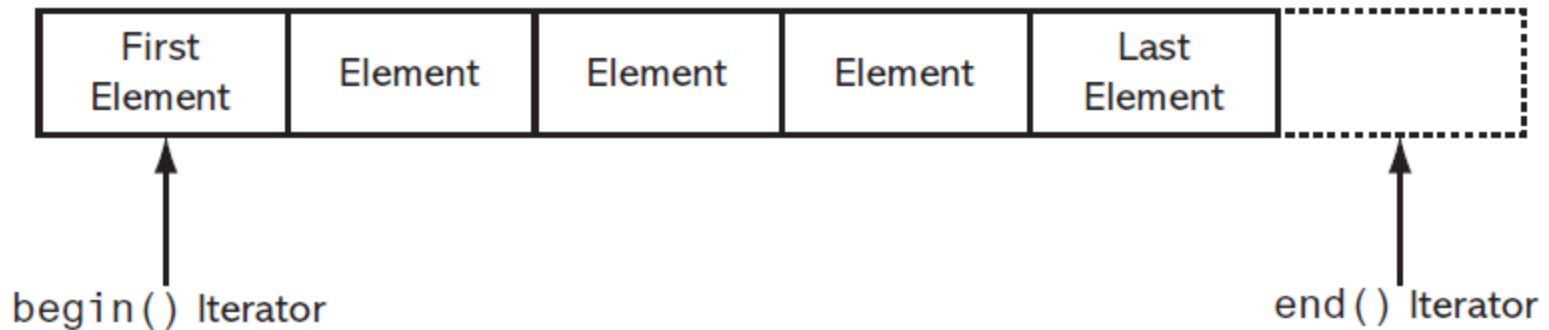
Iterator Examples (2)

- All of the STL containers have a `begin()` member function that returns an iterator pointing to the container's first element.

```
// Define an array object.  
array<string, 3> names = {"Sarah", "William", "Alfredo"};  
  
// Define an iterator for the array object.  
array<string, 3>::iterator it;  
  
// Make the iterator point to the array object's first element.  
it = names.begin();  
  
// Display the element that the iterator points to.  
cout << *it << endl;
```

Iterator Examples (3)

- All of the STL containers have a `end()` member function that returns an iterator pointing to the position *after* the container's last element.



Iterator Examples (4)

- You typically use the `end()` member function to know when you have reached the end of a container.

```
// Define an array object.
array<string, 3> names = {"Sarah", "William", "Alfredo"};

// Define an iterator for the array object.
array<string, 3>::iterator it;

// Make the iterator point to the array object's first element.
it = names.begin();

// Display the array object's contents.
while (it != names.end())
{
    cout << *it << endl;
    it++;
}
```

The vector Class

- A vector is a sequence container that works like an array, but is dynamic in size.
- Overloaded [] operator provides access to existing elements
- The vector class is declared in the `<vector>` header file.

Textbook “Vector-related” Sections

- Section 7.11: Vectors
 - pages 435-449
- Section 8.5: Sorting and Searching Vectors
 - pages 495-498
- Section 17.3: The `vector` Class
 - pages 1040 - 1053

Adding New Elements to a vector

- The `push_back` member function adds a new element to the end of a `vector`:

```
vector<int> numbers;  
numbers.push_back(10);  
numbers.push_back(20);  
numbers.push_back(30);
```

Using an Iterator With a vector

- vectors have `begin()` and `end()` member functions that return iterators pointing to the beginning and end of the container:

```
// Create a vector containing names.
vector<string> names = {"Joe", "Karen", "Lisa", "Jackie"};

// Create an iterator.
vector<string>::iterator it; ← Defines an iterator that is compatible
                               with a vector<string> object

// Use the iterator to display each element in the vector.
for (it = names.begin(); it != names.end(); it++)
{
    cout << *it << endl; ← Displays the item that the iterator points to
}
```

Using an Iterator With a vector

- The `begin()` and `end()` member functions return a random-access iterator of the `iterator` type
- The `cbegin()` and `cend()` member functions return a random-access iterator of the `const_iterator` type
- The `rbegin()` and `rend()` member functions return a reverse iterator of the `reverse_iterator` type
- The `crbegin()` and `crend()` member functions return a reverse iterator of the `const_reverse_iterator` type

The **insert()** Member Function


- You can use the **insert()** member function, along with an iterator, to insert an element at a specific position.
- General format:

vectorName.insert(it, value);

Iterator pointing to an
element in the vector



Value to insert before
the element that *it*
points to



Emplacement of Container Elements (1)

- Member functions such as **insert()** and **push_back()** can cause temporary objects to be created in memory while the insertion is taking place.
- C++11 introduced a new family of member functions that use a technique known as *emplacement* to insert new elements.
- Emplacement avoids the creation of temporary objects in memory while a new object is being inserted into a container.
- The emplacement functions are more efficient than functions such as **insert()** and **push_back()**

Emplacement of Container Elements (2)

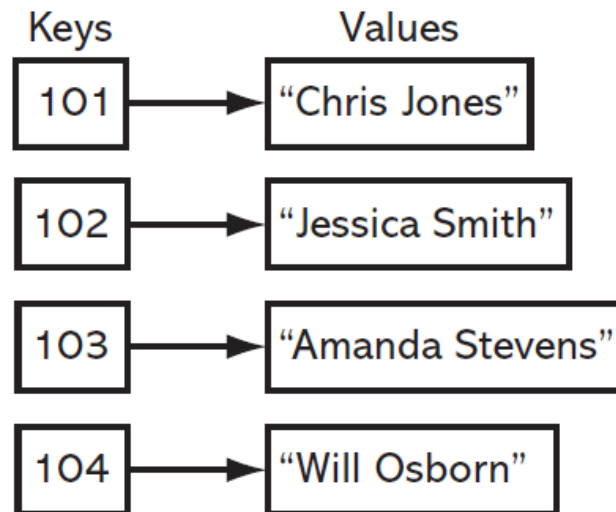
- The `vector` class provides two member functions that use emplacement:
 - **`emplace()`** - emplaces an element at a specific location
 - **`emplace_back()`** - emplaces an element at the end of the vector
- With these member functions, it is not necessary to instantiate, ahead of time, the object you are going to insert.
- Instead, you pass to the emplacement function any arguments that you would normally pass to the constructor of the object you are inserting.
- The emplacement function handles the construction of the object, forwarding the arguments to its constructor.

Maps – General Concepts

- A *map* is an associative container.
- Each element that is stored in a map has two parts: a *key* and a *value*.
- To retrieve a specific value from a map, you use the key that is associated with that value.
- This is similar to the process of looking up a word in the dictionary, where the words are keys and the definitions are values.

Maps

- Example: a map in which employee IDs are the keys and employee names are the values.
- You use an employee's ID to look up that employee's name.



The map Class

- You can use the STL **map** class to store *key-value* pairs.
- The keys that are stored in a **map** container are unique – no duplicates.
- The **map** class is declared in the `<map>` header file.

The map Class -- Example

- Example: defining a map container to hold employee ID numbers (as `ints`) and their corresponding employee names (as `strings`):

```
map<int, string> employees;
```



Key data type Value data type

Initializing a map

```
map<int, string> employees =  
{  
    {101, "Chris Jones"},  
    {102, "Jessica Smith"},  
    {103, "Amanda Stevens"},  
    {104, "Will Osborn"}  
};
```

- In the first element, the key is 101 and the value is "Chris Jones".
- In the second element, the key is 102 and the value is "Jessica Smith".
- In the third element, the key is 103 and the value is "Amanda Stevens".
- In the fourth element, the key is 104 and the value is "Will Osborn".

The Overloaded [] Operator

- You can use the [] operator to add new elements to a map.
- General format:

mapName[key] = value;

- This adds the key-value pair to the map.
- If the key already exists in the map, it's associated value will be changed to *value*.

Deleting Elements

- You can use the `erase()` member function to remove a map element by its key:

```
// Create a map containing employee IDs and names.  
map<int, string> employees =  
{  
    {101, "Chris Jones"}, {102, "Jessica Smith"},  
    {103, "Amanda Stevens"}, {104, "Will Osborn"}  
};
```

```
// Delete the employee with ID 102.  
employees.erase(102);
```



Deletes Jessica Smith from the map

The unordered_map Class


- The unordered_map class is similar to the map class, except in two regards:
 - The keys in an unordered_map are not sorted
 - The unordered_map class has better performance
- You should use the unordered_map class instead of the map class if:
 - You will be making a lot of searches on a large number of elements
 - You are not concerned with retrieving them in key order
- The unordered_map class is declared in the `<unordered_map>` header file

The multimap Class

- The `multimap` class is a map that allows duplicate keys
- The `multimap` class has most of the same member functions as the `map` class (see Table 17-11 in your textbook)
- The `multimap` class is declared in the `<map>` header file

The multimap Class

- Consider a phonebook application where the key is a person's name and the value is that person's phone number.
- A `multimap` container would allow each person to have multiple phone numbers



Sets

- A *set* is an associative container that is similar to a mathematical set.
- You can use the STL `set` class to create a set container.
- All the elements in a `set` must be unique. No two elements can have the same value.
- The elements in a set are automatically sorted in ascending order.
- The `set` class is declared in the `<set>` header file.

The set Class

- You can use the STL set class to create a set container.
- The keys that are stored in a map container are unique – no duplicates.

The set Class

- Example: defining a `set` container to hold integers:

```
set<int> numbers;
```

- Example: defining and initializing a `set` container to hold integers:

```
set<int> numbers = {1, 2, 3, 4, 5};
```

The set Class

- A set cannot contain duplicate items.
- If the same value appears more than once in an initialization list, it will be added to the set only one time.
- For example, the following set will contain the values 1, 2, 3, 4, and 5:

```
set<int> numbers = {1, 1, 2, 2, 3, 4, 5, 5, 5};
```

Adding New Elements to a set

- The `insert()` member function adds a new element to a set:

```
set<int> numbers;  
numbers.insert(10);  
numbers.insert(20);  
numbers.insert(30);
```

Stepping Through a set With the Range-Based for Loop

```
// Create a set containing names.  
set<string> names = {"Joe", "Karen", "Lisa", "Jackie"};  
  
// Display each element.  
for (string element : names)  
{  
    cout << element << endl;  
}
```

Using an Iterator With a set

- The `begin()` and `end()` member functions return a bidirectional iterator of the `iterator` type
- The `cbegin()` and `cend()` member functions return a bidirectional iterator of the `const_iterator` type
- The `rbegin()` and `rend()` member functions return a reverse bidirectional iterator of the `reverse_iterator` type
- The `crbegin()` and `crend()` member functions return a reverse bidirectional iterator of the `const_reverse_iterator` type

STL Algorithms

- The STL provides a number of algorithms, implemented as function templates, in the `<algorithm>` header file.
- These functions perform various operations on ranges of elements.
- A range of elements is a sequence of elements denoted by two iterators:
 - The first iterator points to the first element in the range
 - The second iterator points to the end of the range (the element to which the second iterator points is not included in the range).

Categories of Algorithms in the STL

- Min/max algorithms
- Sorting algorithms
- Search algorithms
- Read-only sequence algorithms
- Copying and moving algorithms
- Swapping algorithms
- Replacement algorithms
- Removal algorithms
- Reversal algorithms
- Fill algorithms
- Rotation algorithms
- Shuffling algorithms
- Set algorithms
- Transformation algorithm
- Partition algorithms
- Merge algorithms
- Permutation algorithms
- Heap algorithms
- Lexicographical comparison algorithm